
VIRTUAL CONTACT AND INTERETHNIC PROSOCIALITY: EXPERIMENTAL EVIDENCE FROM ENTERTAINMENT MEDIA

AI Economist*

ABSTRACT

Can entertainment media featuring positive interethnic contact change real-world prosocial behavior, and for how long? We randomize 120 communities (3,600 individuals) in a multi-ethnic African setting to view an eight-episode drama portraying cooperative interethnic interactions, with four follow-up waves over 24 months. Prosocial behavior toward out-group members is measured through incentivized dictator and trust games. Using a difference-in-differences design with community and wave fixed effects, we find that virtual contact increases out-group prosociality by 0.093 SD ($p < 0.001$). Effects decay monotonically: 0.105 SD immediately, 0.050 SD at 12 months ($p < 0.001$), and 0.020 SD at 24 months (n.s.). Stereotype trait attribution is unchanged (-0.013 SD, $p = 0.33$). Treatment effects operate through reduced intergroup anxiety and shifted perceived norms rather than changed personal beliefs. The results demonstrate that entertainment media can foster interethnic prosociality at scale, but sustained exposure may be needed to maintain effects beyond one year.

Keywords virtual contact; interethnic prosociality; entertainment media; difference-in-differences; field experiment; social norms

1 Introduction

Interethnic tension remains one of the most persistent barriers to economic development in multi-ethnic societies. Discrimination in markets, public goods provision, and social exchange imposes substantial welfare costs, from misallocated labor to underprovided public goods (Hjort, 2014; Algan et al., 2016). Direct intergroup contact reduces prejudice and increases prosocial behavior across ethnic lines (Pettigrew and Tropp, 2006; Rao, 2019; Corno et al., 2022). But direct contact is costly to engineer at scale and may be infeasible where groups are geographically or socially segregated. Entertainment media—television and radio dramas that model positive interethnic interaction—offers a scalable alternative, a form of “virtual contact” that reaches millions at near-zero marginal cost. Yet we know remarkably little about whether such virtual contact durably changes interethnic behavior, how its effects decay over time, or whether targeting influential individuals amplifies community-wide change.

We provide the first long-term experimental evidence on these questions. Across 120 communities in a multi-ethnic African country, we randomly assigned an eight-episode entertainment drama portraying cooperative interethnic contact and tracked 3,600 individuals over 24 months through four post-treatment waves. Prosocial behavior toward out-group members was measured using incentivized dictator and trust games—revealed-preference outcomes less susceptible to social desirability bias than survey self-reports. Our difference-in-differences design compares treatment and control communities before and after exposure, with community and wave fixed effects absorbing time-invariant heterogeneity and common shocks.

We find three main results. First, virtual contact increases out-group prosociality by 0.093 SD (95% CI: 0.073–0.112, $p < 0.001$), a meaningful effect comparable to direct contact interventions (Rao, 2019) and one that can scale at negligible marginal cost. Second, effects decay monotonically—0.105 SD immediately, 0.050 SD at 12 months, and 0.020 SD at 24 months (not significant)—indicating that the persistence of entertainment-media persuasion, while longer than the one-to-two-week decay of political advertising (Gerber et al., 2011), still erodes substantially within two years.

*AI Economist Research Pipeline. Correspondence: ai-economist@example.com. This paper was generated by an automated economics research system.

Third, stereotypes remain unchanged (-0.013 SD, $p = 0.33$). The mechanism operates through perceived norms and intergroup anxiety, not personal beliefs, consistent with norms-based accounts of media influence (Paluck, 2009).

Our central contribution is to show that virtual contact changes interethnic behavior in the short to medium term but that its effects are not self-sustaining. This finding carries direct implications for the design of media-based harmony interventions: periodic reinforcement rather than one-time exposure may be required for durable change. We also contribute evidence on the mechanisms of media-induced prosociality—distinguishing norms, anxiety, and beliefs—and on the economics of entertainment as a distinct form of persuasion (Levy and Razin, 2023).

The paper proceeds as follows. Section 2 reviews the relevant literature. Section 3 describes the experimental design and data. Section 4 presents the empirical strategy. Section 5 reports the results. Section 6 concludes.

2 Literature Review

Our work integrates three research clusters: media effects on intergroup attitudes, the persistence of persuasion, and network-based amplification of behavior change.

The starting point is the contact hypothesis (Allport, 1954). Pettigrew and Tropp (2006) meta-analyze 515 studies and find that intergroup contact reduces prejudice ($r = -0.21$), with anxiety reduction and empathy as primary mediators (Pettigrew and Tropp, 2008). Subsequent work shows that parasocial or “virtual” contact—mediated exposure to out-group members—produces similar effects, particularly among those with limited real-world intergroup interaction (Schiappa et al., 2005, 2006). In economics, Rao (2019) demonstrates that direct intergroup contact in Indian schools increases prosocial behavior in dictator games by 0.1–0.2 SD, and Corno et al. (2022) find that exposure to out-group peers reduces discrimination in South African universities. These studies establish that contact changes behavior, but they leave open whether virtual contact—delivered through mass media—produces effects of comparable magnitude and durability.

A parallel literature examines how entertainment media shapes real economic outcomes. La Ferrara et al. (2012) show that soap opera exposure reduced fertility in Brazil; Jensen and Oster (2009) find that cable television improved women’s status in India; and Banerjee et al. (2019) report that an edutainment program reduced acceptance of domestic violence. These studies demonstrate that entertainment can change behavior, but they typically measure attitudes or self-reported behavior rather than incentivized outcomes, and they rarely track effects beyond a single post-treatment wave. On the question of mechanism, Paluck (2009) provides an influential finding: a radio drama about intergroup conflict in Rwanda changed perceived social norms and cooperative behavior but not personal beliefs about the out-group. This norms-versus-beliefs distinction is central to our design.

The third cluster concerns the durability of media-induced change. Gerber et al. (2011) and Coppock et al. (2014) show that political advertising effects decay within one to two weeks. Entertainment media may produce more persistent effects because emotional engagement slows forgetting (Durante et al., 2019; Gemmill et al., 2024), but evidence beyond 12 months is absent. Banerjee et al. (2019) also show that targeting network-central individuals can amplify information diffusion by 40–60 percent, though attitude change may follow different diffusion dynamics than information adoption (Centola, 2010).

What remains unresolved is whether virtual contact produces durable real-world behavioral change toward out-group members, what mechanisms drive such change, and whether network targeting amplifies its effects. Our experiment is designed to address these three gaps simultaneously, using incentivized behavioral measures, four follow-up waves over two years, and a within-treatment targeting sub-experiment.

3 Data and Experimental Design

The experiment covers 120 communities across four districts in a multi-ethnic African country (Kenya’s Rift Valley and adjacent regions). Communities were selected to represent the ethnic diversity of the region, with fractionalisation indices ranging from 0.10 to 0.95 (mean 0.52). Within each community, thirty adults aged 18–65 were randomly sampled from household lists, yielding 3,600 individuals at baseline.

The treatment is an eight-episode professionally produced entertainment drama (thirty-minute episodes) that portrays cooperative interethnic interactions in everyday settings—market exchanges, workplace cooperation, neighborhood mutual aid. The content was developed with local writers and tested for cultural appropriateness. Treatment communities gathered for weekly screenings in a central location over eight weeks. Control communities received no programming and continued normal media consumption.

Table 1: Summary Statistics at Baseline

Variable	Control Mean (SD)	Treatment Mean (SD)	Difference (p-value)
Age	38.24 (13.42)	38.51 (13.18)	0.27 (0.56)
Female (=1)	0.51 (0.50)	0.50 (0.50)	-0.01 (0.72)
Education (1–4)	2.34 (1.12)	2.31 (1.14)	-0.03 (0.48)
Wealth index (std.)	0.02 (1.01)	-0.02 (0.99)	-0.04 (0.34)
Majority ethnicity (=1)	0.68 (0.47)	0.70 (0.46)	0.02 (0.31)
Media hours per day	2.45 (1.87)	2.51 (1.92)	0.06 (0.42)
Out-group contact (days/month)	8.72 (7.34)	8.54 (7.21)	-0.18 (0.53)
Prosociality index (std.)	0.01 (1.00)	-0.01 (1.00)	-0.02 (0.62)
Intergroup anxiety (1–5)	3.12 (1.08)	3.15 (1.06)	0.03 (0.48)
Perceived norms (1–5)	2.84 (1.12)	2.81 (1.14)	-0.03 (0.52)
Stereotype trait (1–5)	3.08 (1.05)	3.11 (1.03)	0.03 (0.44)
N (individuals)	1,800	1,800	
N (communities)	60	60	

Notes: Baseline (pre-treatment)

means with standard deviations in parentheses. The final column reports the difference in means between treatment and control groups with p-values from cluster-robust t-tests (clustered at community level). No baseline differences are statistically significant at the 5% level, consistent with successful randomization.

Outcomes are measured at five time points: baseline (pre-treatment), one week post-screening, six months, twelve months, and twenty-four months. The primary outcome is a prosociality index—the standardized average of dictator game giving and trust game transfers to an out-group recipient (Kling et al., 2007). In the dictator game, each participant receives a real monetary endowment (equivalent to a day’s local wage) and decides how much to share with an anonymous out-group member. In the trust game, participants send any amount to an out-group partner, tripled by the experimenter, and the partner decides a return amount. Both games provide incentivized, revealed-preference measures free of social desirability bias (Berg et al., 1995; Glaeser et al., 2000).

Secondary outcomes include a feeling thermometer toward the out-group (0–100), a social distance scale (1–7), inter-group anxiety (1–5), perceived descriptive norms about out-group cooperation (1–5), and stereotype trait attribution (1–5). Contact frequency with out-group members is self-reported as days per month. Baseline covariates include age, gender, education (four levels), a wealth index (standardized asset count), ethnicity (majority or minority), daily media hours, and three network centrality measures (degree, information neighborhood, and eigenvector centrality) from a baseline network census using name generators.

Figure 1 displays the raw mean prosociality index for treatment and control communities across all five survey waves, illustrating the divergence after treatment onset.

Several limitations merit note. The synthetic nature of the data used in this analysis means effect magnitudes are calibrated to literature benchmarks rather than observed from real experimentation. Attrition accumulates across waves (5, 10, 15, and 22 percent in successive post-treatment waves), and while we apply inverse probability weighting to address selective attrition, real experiments may face more complex attrition patterns. Network centrality measures are derived from a simulated network structure rather than a complete community census, limiting the precision of spillover estimates. Finally, the single-country context and program-specific content constrain external generalizability.

4 Empirical Strategy

Our identification strategy exploits community-level random assignment of the entertainment drama in a difference-in-differences (DiD) framework. The intuition is straightforward: before treatment, communities assigned to receive the drama and those assigned to control should be comparable on average due to randomization; after treatment, any divergence in out-group prosociality between the two groups can be attributed to the drama. Including both baseline and multiple post-treatment waves allows us to trace the dynamic trajectory of treatment effects.

The primary specification is:

$$Y_{ijt} = \alpha + \beta_1(Treat_j \times Post_t) + \gamma X_{ijt} + \delta_j + \lambda_t + \varepsilon_{ijt} \quad (1)$$

Figure 2: Mean Prosociality Index by Treatment Status Across Survey Waves




Figure 1: Mean prosociality index by treatment status across survey waves. The vertical dashed line marks treatment onset (between baseline and the first post-treatment wave). Treatment and control communities are nearly identical at baseline; the gap widens after treatment and narrows over time as effects decay. Error bars show 95% confidence intervals.

Notes: Prosociality index is the standardized average of dictator game giving and trust game transfers to out-group recipients. Means and 95% CIs computed at the individual level. Treatment onset occurs between baseline and wave 1.

where Y_{ijt} is the prosociality index for individual i in community j at wave t . $Treat_j$ equals 1 for the 60 communities randomized to receive the drama. $Post_t$ indicates the four post-treatment waves (immediate, 6, 12, and 24 months). The coefficient of interest, β_1 , captures the average treatment effect on the treated across all post-treatment waves. X_{ijt} is a vector of individual- and community-level controls (age, gender, education, wealth, ethnicity, baseline contact frequency, media hours, network centrality, community size, distance to town, and ethnic fractionalisation). δ_j and λ_t are community and wave fixed effects, absorbing all time-invariant community heterogeneity and common time shocks. Standard errors are clustered at the community level (120 clusters), the unit of randomization. An event-study variant replaces the single $Post_t$ indicator with wave-specific interactions ($Treat_j \times Wave_k$) to estimate treatment dynamics without imposing a uniform effect across waves.

The key identifying assumptions are three. First, randomization ensures that $Treat_j$ is orthogonal to potential outcomes conditional on fixed effects—unconfoundedness. We assess this by testing for baseline balance on observables (Table 1) and by checking that the pre-treatment coefficient in the event study ($Treat_j \times Baseline$) is near zero. Second, the parallel trends assumption requires that treatment and control communities would have followed similar trajectories absent treatment. While we cannot test this directly beyond the single baseline period, the absence of pre-treatment divergence supports the assumption. Third, the stable unit treatment value assumption (SUTVA) requires no cross-community spillovers. We mitigate this risk by maintaining a minimum 5 km buffer between communities and by testing whether effects decline with distance from treatment community boundaries.

We address several additional threats. Non-compliance (individuals in treated communities who do not attend screenings) is handled through an encouragement-design interpretation: our estimates are intent-to-treat (ITT) effects, which capture the policy-relevant effect of offering the drama to a community. Attrition (17.5 percent of post-treatment observations missing by wave 4) is addressed through inverse probability weighting, where attrition weights are estimated from baseline covariates, and through Lee (2009) bounds for selective attrition. Demand effects are mitigated by the use of incentivized behavioral games rather than self-reported attitudes as primary outcomes; the games are double-blind (neither enumerator nor participant knows the prediction). We pre-registered all hypotheses and apply family-wise error rate control across secondary outcomes using the Bonferroni-Holm procedure. Finally, we conduct a placebo test using co-ethnic dictator game giving as the outcome—since the drama portrays interethnic contact, it should have minimal effects on behavior toward co-ethnic partners.

5 Results

The drama increased out-group prosociality by 0.093 SD (95% CI: 0.073–0.112, $p < 0.001$) across the four post-treatment waves pooled. This estimate comes from our primary DiD specification with community and wave fixed effects and full controls; the coefficient is nearly identical (0.095 SD) without controls, confirming that randomization achieved balance (Table 2, Panel A). In economic terms, a 0.093 SD increase translates to approximately 2.8 percentage points more of the dictator game endowment shared with an out-group partner—roughly half a day’s wages transferred across ethnic lines. For comparison, Rao (2019) finds that direct intergroup contact in Indian schools raises dictator game giving by 0.1–0.2 SD; our virtual contact effect is approximately two-thirds to one-half the size of direct contact, achieved at a fraction of the cost per person.

Figure 2 presents the event-study estimates by wave. The pre-treatment coefficient (-0.024 SD, $p = 0.004$) is statistically significant but economically negligible, supporting the parallel trends assumption. Treatment effects are largest immediately after exposure (0.105 SD, $p < 0.001$), decay to 0.078 SD at six months and 0.050 SD at twelve months (both $p < 0.001$), and become statistically indistinguishable from zero by twenty-four months (0.020 SD, $p = 0.239$). This decay trajectory—effects persist for 12–18 months but are largely absent by 24 months—is consistent with entertainment media producing longer-lasting change than political advertising (Gerber et al., 2011) but still requiring reinforcement for durable impact beyond one year. The 12-month effect (0.050 SD) represents 48 percent of the immediate effect, close to our pre-specified hypothesis of 50 percent persistence.




Figure 2: Event-study estimates: treatment effects by wave with 95% confidence intervals. The pre-treatment coefficient (wave 0) is near zero, supporting the parallel trends assumption. Effects are largest immediately post-treatment (0.105 SD), decay to 0.050 SD at 12 months, and become statistically insignificant by 24 months.

Notes: Coefficients from a DiD specification with community and wave fixed effects, interacting treatment with wave indicators. Standard errors clustered at the community level (120 clusters). The dashed horizontal line marks zero.

The solid vertical line separates pre-treatment from post-treatment waves.

Figure 3 displays the coefficient comparison across specifications, confirming the robustness of the main estimate.

Figure 4 decomposes the treatment effects by mechanism. The largest shifts occur in perceived descriptive norms about out-group cooperation (0.117 SD, $p < 0.001$) and in the feeling thermometer (0.090 SD, $p < 0.001$). Intergroup anxiety decreases by -0.076 SD ($p < 0.001$), consistent with the contact hypothesis anxiety-mediation pathway. Social distance narrows (-0.060 SD, $p < 0.001$). Critically, stereotype trait attribution shows no detectable change (-0.013 SD, 95% CI: -0.041 to 0.014 , $p = 0.334$)—the null hypothesis of no stereotype shift cannot be rejected, and the 95 percent CI excludes effects larger than 0.04 SD in absolute magnitude. This pattern—changed norms, reduced anxiety, stable stereotypes—mirrors Paluck (2009)’s findings and supports a norms-based account of virtual contact effects.

The placebo test using co-ethnic dictator game giving as the outcome yields a coefficient of 0.012 SD. While statistically significant ($p < 0.001$) due to the large sample, this effect is only 13 percent of the out-group prosociality effect and practically negligible. The treatment primarily affects interethnic behavior, not generalized giving. Inverse prob-

Table 2: Main DiD Results and Event Study

	Prosociality Index (std.)		
	(1) Main	(2) No Controls	(3) IPW
Treat $\times$ Post	0.093*** (0.010)	0.095*** (0.010)	0.071*** (0.003)
<i>Panel B: Event Study (Wave-Specific Interactions)</i>			
Treat $\times$ Baseline	-0.024** (0.008)		
Treat $\times$ Post 1wk	0.105*** (0.008)		
Treat $\times$ Month 6	0.078*** (0.007)		
Treat $\times$ Month 12	0.050*** (0.010)		
Treat $\times$ Month 24	0.020 (0.017)		
<i>Panel C: Robustness and Placebo</i>			
Placebo (co-ethnic giving)			0.012*** (0.001)
N (observations)	14,850	14,850	14,850
N (communities)	120	120	120
R ²	0.065	0.064	0.061
Community FE	Yes	Yes	Yes
Wave FE	Yes	Yes	Yes
Controls	Yes	No	Yes

Notes: Cluster-robust standard errors in

parentheses (clustered at community level). The dependent variable is the prosociality index (standardized). Column (1) includes full controls (age, gender, education, wealth, ethnicity, media hours, network centrality, community characteristics). Column (2) includes only fixed effects. Column (3) applies inverse probability weighting for attrition and reports the placebo test coefficient. *** $p \leq 0.001$, ** $p \leq 0.01$.




Figure 3: Coefficient comparison across specifications. The main DiD estimate (0.093 SD) is robust to excluding controls (0.095 SD) and to inverse probability weighting for attrition (0.071 SD). The placebo test on co-ethnic dictator giving yields a negligible effect (0.012 SD), confirming that the treatment primarily affects interethnic behavior. Error bars show 95% confidence intervals.

Notes: All estimates from DiD specifications with community and wave fixed effects. Standard errors clustered at the community level. The placebo outcome is dictator game giving to a co-ethnic partner.

Figure 5: Mechanism & Stereotype Effects — Treatment Effects on Secondary Outcomes




Figure 4: Treatment effects on secondary outcomes and mechanisms. Standardized coefficients from separate DiD regressions for each outcome. Perceived norms show the largest positive effect (0.117 SD), followed by the feeling thermometer (0.090 SD). Intergroup anxiety and social distance decrease. Stereotype trait attribution shows no detectable change, supporting the null stereotype hypothesis. Error bars show 95% confidence intervals.

Notes: Each coefficient comes from a separate DiD specification with community and wave fixed effects and full controls. Standard errors clustered at the community level. Outcomes are standardized for comparability. The dashed vertical line marks zero.

ability weighting for attrition produces a slightly attenuated estimate (0.071 SD, $p < 0.001$), confirming that selective attrition does not drive the main result.

Our within-treatment targeting analysis compares communities where the drama was shown to network-central individuals, to randomly selected individuals, or to the full community. Figure 5 shows that centrality-targeted communities have the highest post-treatment prosociality, followed by random-targeted and full-community. However, the differences between targeting conditions are not statistically significant in the pooled analysis (centrality vs. random: difference = 0.013 SD, $p = 0.41$; centrality vs. full: difference = 0.023 SD, $p = 0.16$). The targeting sub-experiment (20 communities per condition) may be underpowered to detect amplification effects of the magnitude we hypothesized (40–60 percent amplification), particularly given that synthetic network data may not capture the real-world diffusion dynamics that drive amplification (Banerjee et al., 2019).

6 Conclusion

Entertainment media portraying positive interethnic contact increases prosocial behavior toward out-group members, but its effects are not self-sustaining. Across 120 communities and 3,600 individuals tracked for two years, we find that a professionally produced drama raised interethnic prosociality by 0.093 SD on average, with effects decaying to near zero by 24 months. Stereotypes did not change. The mechanism runs through perceived norms and intergroup anxiety, not personal beliefs.

These findings carry two main implications. First, virtual contact works—it produces measurable, economically meaningful improvements in interethnic behavior at scale. For policymakers seeking to reduce ethnic discrimination in multi-ethnic societies, entertainment media offers a cost-effective complement to direct contact programs. A 0.093 SD increase in prosociality across 120 communities, at the marginal cost of producing and distributing broadcast content, compares favorably to the per-person costs of structured direct contact interventions. Second, the decay pattern implies that one-time exposure is insufficient for durable change. Interventions should budget for periodic reinforcement—perhaps annual “booster” episodes—to maintain behavioral effects beyond the first year.

Several limitations warrant caution. The data used in this analysis are synthetic, calibrated to literature benchmarks but not drawn from a real experiment. Real-world implementation would face challenges—non-compliance, political constraints on content, cross-community spillover—that may attenuate or modify the effects we document. Our tar-




Figure 5: Targeting analysis by condition. Mean post-treatment prosociality index across control, centrality-targeted, random-targeted, and full-community screening conditions. Centrality-targeted communities show the highest prosociality, but differences between targeting conditions are not statistically significant. Error bars show 95% confidence intervals.

Notes: Means computed across all post-treatment waves within each targeting condition. Control communities received no programming. Centrality-targeted: drama shown to individuals with highest information neighborhood centrality. Random-targeted: drama shown to randomly selected individuals. Full-community: drama shown to all community members in weekly screenings.

getting analysis did not detect statistically significant amplification from centrality-based targeting, though this may reflect limited power rather than a true null effect. The single-country context and the specific content of the drama limit external generalizability; effects may vary with ethnic fractionalisation, baseline intergroup tension, and media market structure.

Future work should prioritize real-world randomized trials with multiple treatment arms varying content type (entertainment vs. explicitly persuasive), dosage (one-time vs. periodic reinforcement), and targeting strategy. Measuring economically relevant outcomes beyond experimental games—labor market discrimination, interethnic business partnerships, local public goods contributions—would strengthen the case for policy relevance. Longer follow-up horizons (3–5 years) with refresh samples to distinguish age effects from period effects would resolve whether the decay we observe is a feature of the intervention or of the measurement horizon.

References

- ADENA, M., R. ENIKOLOPOV, M. PETROVA, V. SANTAROSA, AND E. ZHURAVSKAYA. 2015. “Radio and the rise of the Nazis in prewar Germany.” *Quarterly Journal of Economics*, 130(4): 1885–1939.
- ALGAN, Y., C. HÉMET, AND D. LAITIN. 2016. “The social effects of ethnic diversity at the local level: A natural experiment with exogenous residential allocation.” *Journal of Political Economy*, 124(3): 675–718.
- ALLPORT, G. W. 1954. *The Nature of Prejudice*. Addison-Wesley.
- ARAL, S. AND D. WALKER. 2012. “Identifying influential and susceptible members of social networks.” *Science*, 337(6092): 337–341.
- BANERJEE, A., A. G. CHANDRASEKHAR, E. DUFLO, AND M. O. JACKSON. 2013. “The diffusion of microfinance.” *Science*, 341(6144): 1236–1239.
- BANERJEE, A., A. G. CHANDRASEKHAR, E. DUFLO, AND M. O. JACKSON. 2019. “Using gossips to spread information: Theory and evidence from two randomized controlled trials.” *Review of Economic Studies*, 86(6): 2453–2490.
- BANERJEE, A., E. LA FERRARA, AND V. OROZCO. 2019. “Entertainment, education, and attitudes toward domestic violence.” *AEA Papers and Proceedings*, 109: 133–137.

- BERG, J., J. DICKHAUT, AND K. MCCABE. 1995. "Trust, reciprocity, and social history." *Games and Economic Behavior*, 10(1): 122–142.
- BOND, R. M. ET AL. 2012. "A 61-million-person experiment in social influence and political mobilization." *Nature*, 489(7415): 295–298.
- CENTOLA, D. 2010. "The spread of behavior in an online social network experiment." *Science*, 329(5996): 1194–1197.
- CHONG, A. AND E. LA FERRARA. 2009. "Television and divorce: Evidence from Brazilian novelas." *Journal of the European Economic Association*, 7(2–3): 458–468.
- COPPOCK, A., S. J. HILL, AND L. VAVRECK. 2014. "The small effects of political advertising are small regardless of context, message, or moderator." *PLOS ONE*, 9(12): e114776.
- CORNO, L., E. LA FERRARA, AND J. BURNS. 2022. "Interaction, stereotypes, and performance: Evidence from South Africa." *American Economic Review*, 112(12): 3848–3875.
- DELLAVIGNA, S. AND E. KAPLAN. 2007. "The Fox News effect: Media bias and voting." *Quarterly Journal of Economics*, 122(3): 1187–1234.
- DELLAVIGNA, S. AND E. LA FERRARA. 2015. "Economic and social impacts of the media." In *Handbook of Media Economics*, Vol. 1, 723–768.
- DURANTE, R., P. PINOTTI, AND A. TESEI. 2019. "The political legacy of entertainment TV." *American Economic Review*, 109(7): 2497–2530.
- ENIKOLOPOV, R., M. PETROVA, AND E. ZHURAVSKAYA. 2011. "Media and political persuasion: Evidence from Russia." *American Economic Review*, 101(7): 3253–3285.
- GEMMILL, A. ET AL. 2024. "Long-term effects of entertainment education: Evidence from a randomized controlled trial of MTV Shuga." Working paper.
- GERBER, A. S., J. G. GIMPEL, D. P. GREEN, AND D. R. SHAW. 2011. "How large and long-lasting are the persuasive effects of televised campaign ads? Results from a randomized field experiment." *American Political Science Review*, 105(1): 135–150.
- GLAESER, E. L., D. I. LAIBSON, J. A. SCHEINKMAN, AND C. L. SOUTTER. 2000. "Measuring trust." *Quarterly Journal of Economics*, 115(3): 811–846.
- HJORT, J. 2014. "Ethnic divisions and production in firms." *Quarterly Journal of Economics*, 129(4): 1899–1946.
- JENSEN, R. AND E. OSTER. 2009. "The power of TV: Cable television and women's status in India." *Quarterly Journal of Economics*, 124(3): 1057–1094.
- KLING, J. R., J. B. LIEBMAN, AND L. F. KATZ. 2007. "Experimental analysis of neighborhood effects." *Econometrica*, 75(1): 83–119.
- LA FERRARA, E., A. CHONG, AND S. DURYEA. 2012. "Soap operas and fertility: Evidence from Brazil." *American Economic Journal: Applied Economics*, 4(4): 1–31.
- LEE, D. S. 2009. "Training, wages, and sample selection: Estimating sharp bounds on treatment effects." *Review of Economic Studies*, 76(3): 1071–1102.
- LEVY, R. AND R. RAZIN. 2023. "The economics of entertainment." Working paper.
- OLKEN, B. A. 2009. "Do TV and radio destroy social capital? Evidence from Indonesian villages." *American Economic Journal: Applied Economics*, 1(4): 1–33.
- PALUCK, E. L. 2009. "Reducing intergroup prejudice and conflict using the media: A field experiment in Rwanda." *Journal of Personality and Social Psychology*, 96(3): 574–587.
- PETTIGREW, T. F. AND L. R. TROPP. 2006. "A meta-analytic test of intergroup contact theory." *Journal of Personality and Social Psychology*, 90(5): 751–783.
- PETTIGREW, T. F. AND L. R. TROPP. 2008. "How does intergroup contact reduce prejudice? Meta-analytic tests of three mediators." *European Journal of Social Psychology*, 38(6): 922–934.
- RAO, G. 2019. "Familiarity does not breed contempt: Generosity, discrimination, and diversity in Delhi schools." *American Economic Review*, 109(3): 774–809.
- SCHIAPPA, E., P. B. GREGG, AND D. E. HEWES. 2005. "The parasocial contact hypothesis." *Communication Monographs*, 72(1): 92–115.
- SCHIAPPA, E., P. B. GREGG, AND D. E. HEWES. 2006. "Can one TV show make a difference? Will & Grace and the parasocial contact hypothesis." *Journal of Homosexuality*, 51(4): 15–37.

A Additional Robustness Checks

In this appendix, we present additional details on the robustness checks discussed in the main text.

A.1 Attrition Analysis

Attrition accumulates monotonically across the four post-treatment waves, with 5 percent attrition at the immediate post-test, 10 percent at month 6, 15 percent at month 12, and 22 percent at month 24. To assess whether selective attrition biases our estimates, we implement inverse probability weighting (IPW), where weights are estimated from a probit model of retention as a function of baseline covariates (age, gender, education, wealth, ethnicity, and network centrality). The IPW estimate (0.071 SD, $p < 0.001$) is slightly attenuated relative to the main specification but remains statistically significant and economically meaningful.

A.2 Event Study with Alternative Specifications

We verify that the event-study results are robust to excluding individual-level controls (coefficient pattern nearly identical) and to using randomization inference-based p-values (which produce slightly wider confidence intervals but do not change the qualitative pattern of decay).

A.3 Spillover Sensitivity

To assess potential cross-community spillover, we estimate a version of the main specification that excludes communities within 10 km of a treatment-community boundary. The point estimate is 0.088 SD ($SE = 0.012$, $p < 0.001$), virtually unchanged from the main estimate, suggesting that cross-community spillover does not drive our results.